

ANNOT8

Conversational PDF Annotation with Large Language Models


Paul Fishwick and Claude Code

May 25, 2026

`https://metaphorz.github.io/annot8/`

The Idea

- Technical PDFs — patents, standards, reports — are long and dense.
- Readers want to **ask questions** and keep the answers **attached to the source text**.
- **ANNOT8**: converse with an LLM about a PDF; every answer becomes an annotation on the passage it draws from.
- Motivated by a 304-page *scanned* patent.




Answers auto-annotated, colour per question.

How It Works

- Drag-drop a PDF — it stays in your browser.
- Ask about the **whole document** or the current page.
- Each answer is **auto-annotated**; the highlight colour cycles per question.
- Export a PDF with real Acrobat-style annotations.

Two ways to reach an LLM:

- **Copy/paste** (default) — use your own LLM.
- **OpenRouter API** — latest Claude, GPT, Gemini.



Copy/paste mode: ANNOT8 builds the prompt; the reply is pasted back and applied locally.

- **Text extraction** — native PDF text layer per page, with Tesseract OCR fallback for scanned pages.
- **Fuzzy quote location** — the hard part:
 - The LLM's page numbers *drift* — off by 39–79 pages on the patent.
 - OCR text is never copied verbatim.
 - Content-token windowed matching searches *every* page and trims the highlight to the densest cluster.
- **Ten** Adobe Acrobat-style annotation tools; six colours.
- Two builds: a local Flask app, and a **100% browser-only** build hosted on GitHub Pages.

Security — No Data Leak

For an enterprise user on a private, approved LLM, **nothing the user enters may leak.**

- Copy/paste mode: the prompt is shown for *you* to copy — ANNOT8 never transmits it.
- All libraries *and* the Tesseract OCR engine are vendored — works offline and behind corporate proxies.
- Self-hostable, so even page load never contacts GitHub.

Session (network capture)	Requests	External
Copy/paste: load → annotate → export	8 (own host)	none
Scanned-PDF load + in-browser OCR	10 (own host)	none

Results

- Automated tests: backend **29/29**, local UI **9/9**, browser build **10/10** (Selenium).
- Whole-document summary of the 304-page patent: **5/5** key points located and highlighted, despite page drift up to 79 pages.
- Annotations export as real objects — visible in Acrobat's Comments panel.

[Try it](#) / [read more](#)

App: <https://metaphorz.github.io/annot8/>

Report: <https://metaphorz.github.io/annot8/report/annot8-report.pdf>

Source: <https://github.com/metaphorz/annot8>

ANNOT8 — crafted by Paul Fishwick and Claude Code.